

compiled Project

Seth Billiau

4/30/2020

EDA

Read in the data and make note of missing values:

```
data_raw = read.csv("data/chd_risk.csv")
summary(data_raw)
```

```
##      age                education      cigsPerDay      totChol
##  Min.   :32.00  College or Higher : 473  Min.    : 0.000  Min.    :107.0
##  1st Qu.:42.00  High School or GED:1253  1st Qu.: 0.000  1st Qu.:206.0
##  Median :49.00  Some College       : 687  Median : 0.000  Median :234.0
##  Mean   :49.58  Some High School   :1720  Mean    : 9.003  Mean    :236.7
##  3rd Qu.:56.00  NA's               : 105  3rd Qu.:20.000  3rd Qu.:263.0
##  Max.   :70.00                      Max.    :70.000  Max.    :696.0
##                                NA's    :29      NA's    :50
##      sysBP      diaBP      BMI      heartRate
##  Min.   : 83.5  Min.    : 48.00  Min.    :15.54  Min.    : 44.00
##  1st Qu.:117.0  1st Qu.: 75.00  1st Qu.:23.07  1st Qu.: 68.00
##  Median :128.0  Median : 82.00  Median :25.40  Median : 75.00
##  Mean   :132.4  Mean    : 82.89  Mean    :25.80  Mean    : 75.88
##  3rd Qu.:144.0  3rd Qu.: 89.88  3rd Qu.:28.04  3rd Qu.: 83.00
##  Max.   :295.0  Max.    :142.50  Max.    :56.80  Max.    :143.00
##                                NA's    :19      NA's    :1
##      glucose      sex      smoker      OnBPMeds      PrevStroke
##  Min.    : 40.00  female:2419  Nonsmoker:2144  No :4061  No :4213
##  1st Qu.: 71.00  male :1819  Smoker :2094   Yes : 124  Yes:  25
##  Median : 78.00                                NA's:  53
##  Mean    : 81.97
##  3rd Qu.: 87.00
##  Max.    :394.00
##  NA's    :388
##      Hyp      Diab      CHD_Risk
##  No :2922  No :4129  No :3594
##  Yes:1316  Yes: 109  Yes: 644
##
##
##
##
##
```

Check for missingness

Count number of missing predictors in each variable:

```
# Generate the number of missing values for each predictor  
apply(is.na(data_raw), 2, sum)
```

```
##      age  education  cigsPerDay   totChol   sysBP   diaBP   BMI  
##      0      105      29         50        0        0      19  
## heartRate  glucose      sex   smoker  OnBPMeds PrevStroke  Hyp  
##      1      388        0        0        53        0      0  
##      Diab  CHD_Risk  
##      0        0
```

```
missing_preds = c("education", "cigsPerDay", "totChol", "BMI",  
                  "heartRate", "glucose", "OnBPMeds")
```

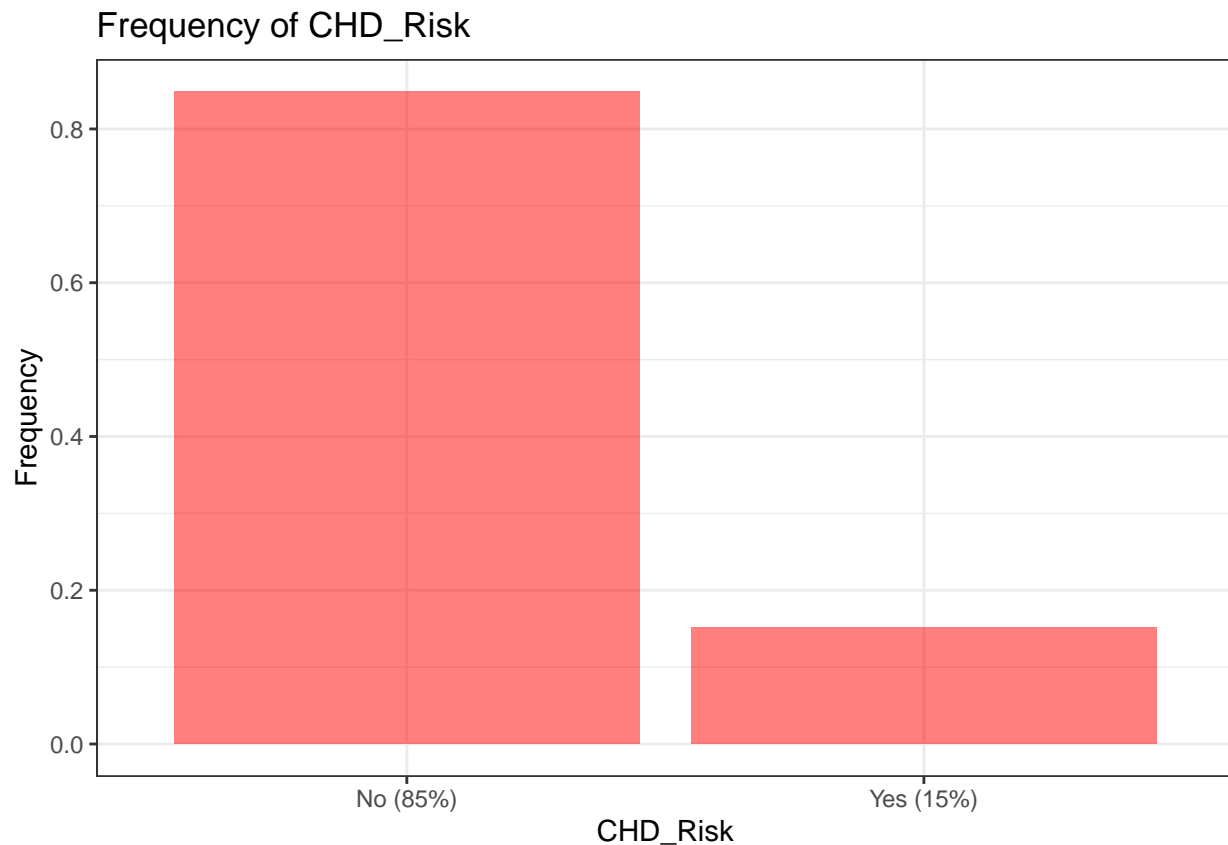
Visualizing the Response

```
# Labels for ticks  
N_label = paste("No (", round(mean(data_raw$CHD_Risk == "No"), 2)*100, "%)", sep="")  
Y_label = paste("Yes (", round(mean(data_raw$CHD_Risk == "Yes"), 2)*100, "%)", sep="")
```

```
# Plot bar chart  
p <- ggplot(data=data_raw, aes(x=as.factor(CHD_Risk))) +  
  geom_histogram(stat="count", fill="red",  
                aes(y=..count../sum(..count..)),  
                alpha = 0.5) +  
  labs(title="Frequency of CHD_Risk") +  
  xlab("CHD_Risk") +  
  ylab("Frequency") +  
  scale_x_discrete(breaks=c("No", "Yes"), labels=c(N_label, Y_label)) +  
  theme_bw()
```

```
## Warning: Ignoring unknown parameters: binwidth, bins, pad
```

```
p
```



Visualizing the Predictors

Visualize distribution of quantitative predictors conditional on the CHD outcome:

```
quant_preds = c("age", "cigsPerDay", "totChol", "sysBP",
                "diaBP", "BMI", "heartRate", "glucose")

make_cond_hist = function(varname) {
  p1 = ggplot(data_raw, aes_string(x=varname)) +
    geom_histogram(aes(y = ..density..),
                  fill = "red", alpha = 0.5) +
    labs(title=paste(varname, "given CHD_Risk")) +
    xlab(varname) +
    ylab("Density") +
    facet_grid(. ~ CHD_Risk) +
    theme_bw()
  return(p1)
}

graphs = lapply(quant_preds, make_cond_hist)
figure1 = ggarrange(graphs[[1]], graphs[[2]], graphs[[3]], graphs[[4]],
                    graphs[[5]], graphs[[6]], graphs[[7]], graphs[[8]],
                    ncol = 2, nrow = 4)
annotate_figure(figure1,
                top = text_grob("Visualizing Quantitative Predictors given CHD_Risk (prevalence = 0.152)"))
```

Visualizing Quantitative Predictors given CHD_Risk (prevalence = 0.152)

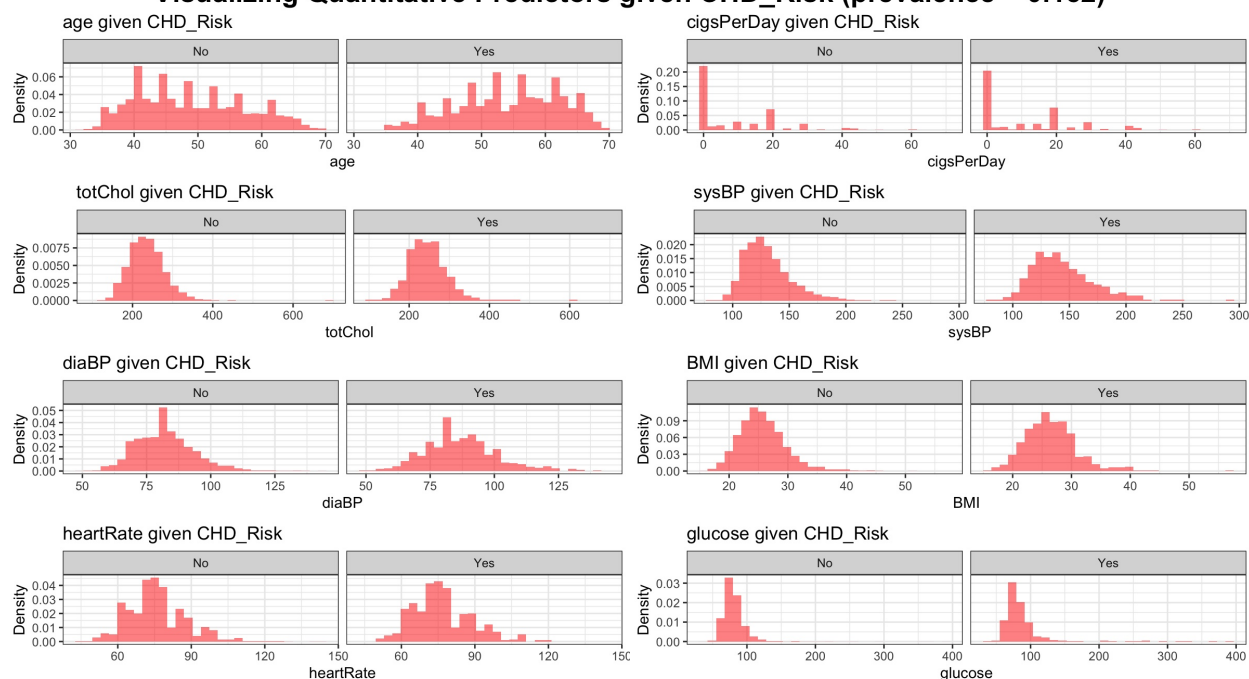


Figure 1: Quantitative EDA

Visualizing the Qualitative predictors by showing their distributions conditional on the outcome:

```
# Address Categorical predictors
cat_preds = c("education", "sex", "smoker", "OnBPMeds",
              "PrevStroke", "Hyp", "Diab")

get_cond_prob_table = function(TABLE) {
  col1 = TABLE[,1] / sum(TABLE[,1])
  col2 = TABLE[,2] / sum(TABLE[,2])
  return(cbind(No=col1, Yes=col2))
}

tab_education = get_cond_prob_table(table(data_raw$education, data_raw$CHD_Risk))
tab_sex = get_cond_prob_table(table(data_raw$sex, data_raw$CHD_Risk))
tab_smoker = get_cond_prob_table(table(data_raw$smoker, data_raw$CHD_Risk))
tab_OnBPMeds = get_cond_prob_table(table(data_raw$OnBPMeds, data_raw$CHD_Risk))
tab_PrevStroke = get_cond_prob_table(table(data_raw$PrevStroke, data_raw$CHD_Risk))
tab_Hyp = get_cond_prob_table(table(data_raw$Hyp, data_raw$CHD_Risk))
tab_Diab = get_cond_prob_table(table(data_raw$Diab, data_raw$CHD_Risk))

tab_prob_Yes = rbind(tab_education, tab_sex, tab_smoker,
                    tab_OnBPMeds, tab_PrevStroke, tab_Hyp,
                    tab_Diab)
round(tab_prob_Yes,3)
```

```
##           No    Yes
## College or Higher 0.115 0.111
## High School or GED 0.316 0.234
## Some College      0.171 0.140
## Some High School  0.399 0.514
## female            0.589 0.467
```

```
## male          0.411 0.533
## Nonsmoker     0.510 0.483
## Smoker        0.490 0.517
## No            0.977 0.935
## Yes           0.023 0.065
## No            0.996 0.983
## Yes           0.004 0.017
## No            0.724 0.495
## Yes           0.276 0.505
## No            0.981 0.938
## Yes           0.019 0.062
```

Check for collinearity with GVIF.

```
# Check for multicollinearity
mod.vif.lm <- lm(as.numeric(CHD_Risk) ~ ., data=data_raw)
vif(mod.vif.lm)
```

```
##          GVIF Df GVIF^(1/(2*Df))
## age      1.397737 1      1.182259
## education 1.124453 3      1.019742
## cigsPerDay 2.732416 1      1.653002
## totChol   1.116842 1      1.056808
## sysBP     3.767158 1      1.940917
## diaBP     3.000260 1      1.732126
## BMI       1.246685 1      1.116550
## heartRate 1.095015 1      1.046429
## glucose   1.638312 1      1.279966
## sex       1.223718 1      1.106218
## smoker    2.585357 1      1.607904
## OnBPMeds  1.111774 1      1.054407
## PrevStroke 1.017647 1      1.008785
## Hyp       2.051447 1      1.432287
## Diab      1.616622 1      1.271465
```

Because all values in the last column are less than $3.1623 = \sqrt{10}$, there is not significant/strong evidence of multicollinearity.